

Wonders Kingdom 11.0 vs 10.2, a build quality read

Prepared for Lea in Publishing and Marc in Game Design. Android, worldwide. Extract of May 3, 2026. Both builds are compared on the matched install window April 15 to April 26, 2026, where they run side by side and every cohort is mature to day 7.

Hold the 100% rollout. 11.0 is worse than 10.2.

Retention is lower at every horizon and the gap widens with cohort age, so this is structural, not a launch blip. The cause is concentrated. Level 7 broke in 11.0. Fix it, verify on a fresh cohort, then resume the ramp.

50.1%

Day 1 retention on 11.0
10% below 10.2

12.9%

Day 7 retention on 11.0
34% below 10.2

39.6%


Level 7 completion on 11.0
was 85.4% on 10.2

31 min

Daily play per player on 11.0
was 44 min on 10.2

1 Retention, for Lea

outputs/metrics.json · retention = active users / day 0 active users · cohort size 22,334 (10.2) vs 10,595 (11.0)




Horizon	10.2	11.0	Gap	vs 10.2
Day 1	55.4%	50.1%	5.4 pts	10% lower
Day 3	31.9%	24.2%	7.6 pts	24% lower
Day 7	19.6%	12.9%	6.7 pts	34% lower

11.0 keeps shedding the players that 10.2 keeps. The gap opens at day 1 and grows to day 7, and it is not noise. Day 1 is lower on every one of the 12 shared install dates. Even the players who stay play less, from 44 to 31 minutes a day and from 2.64 to 2.23 sessions. The US matches the world, day 7 reads 12.9 vs 19.4, so this is the build and not a region.

2 What broke, for Marc

outputs/metrics.json · completion = completed / started · ratio metrics, so retries never inflate a level



Level	Completion 10.2	Completion 11.0	Change	Attempts per user	Seconds per attempt
Level 6	94.1%	92.3%	2 pts lower	1.05 to 1.07	160 to 159
Level 7	85.4%	39.6%	46 pts lower	1.14 to 1.94	159 to 162
Level 8	82.1%	74.1%	8 pts lower	1.17 to 1.27	160 to 157

Level 7 is the wall. Completion drops 46 points and retries almost double. It breaks in all 10 countries and from the very first 11.0 cohort, so it shipped broken. Seconds per attempt hold at 162, so attempts run their full course and fail. This reads like a difficulty or win condition change, not a crash. The wall costs real players. The share reaching level 8 falls from 42.7% to 31.7%.

A milder drift sits under the break. Levels 1 to 6 average 2.5 points lower and levels 8 to 24 average 8.1 points lower, with retries up across the curve. Fix level 7 first, then review the wider difficulty pass.

3 Why you can trust this read

the three choices that carry the verdict

Naive pooled read, day 7 on 11.0



6.8%

Corrected read, matched window



12.9%

Right censoring

The matched window removes it

Young 11.0 cohorts have no day 7 row yet, so a naive pool reports 6.8% and nearly doubles the apparent damage. The matched window and a full history maturity gate both return 12.9%.

Version labels

The 11.0.1 label is quarantined

Its installs start on February 2, 2026, two months before 11.0 exists, which is impossible for a patch. It is 767 players, about 0.3% of volume. It behaves like 11.0, so it is most likely mistagged. It is flagged to data engineering and kept out of the headline.

Rollout mix

The comparison is clean

Country, network and package shares match within 0.7 points between the builds. The pooled worldwide read is valid and needs no reweighting.

Recommendation

Fix level 7 in 11.0 first. Roll back or rebalance the change. Confirm completion returns near 85% and that day 1 and day 3 recover on a fresh cohort. Then resume the rollout. Investigate the wider difficulty drift across the rest of the curve next.

4 Notes for the reviewer

the two the brief asks for, plus the production guard, in a line each

Modelling choice

One window, one quarantine

Compare the builds only on the matched install window April 15 to 26, where both run side by side and every cohort is mature to day 7, and quarantine the impossible 11.0.1 label. That single choice removes right censoring, calendar mismatch and label noise at once, and it moves 11.0 day 7 from a censored 6.8% to the real 12.9%.

LLM tools

Scaffolding, checked at every number

Claude Code scaffolded the DuckDB pipeline, the SQL and the charts. Its first retention pass pooled every 11.0 cohort and reported day 7 at 6.8%, which halves the real 12.9% because young cohorts have no day 7 row yet. I caught it as right censoring, and also blocked a funnel built on `i_level_started`, which counts retries and would inflate the very level that broke.

If this ran nightly, key the maturity gate off the run date, sanity check every new version label against its release date, and keep re aggregation idempotent for the latest, still arriving day. Otherwise every fresh build fakes a day 7 collapse.